

Graph algorithms - Bellman-Ford algorithm

Problem

Given a graph with no negative costs cycles and two vertices s and t , find a minimum cost walk from s to t .

Idea

The algorithm keeps two mappings:

- $dist[x]$ = the cost of the minimum cost walk from s to x known so far
- $prev[x]$ = the vertex just before x on the walk above.

Initially, $dist[s]=0$ and $dist[x]=\infty$ for $x \neq s$; this reflects the fact that we only know a zero-length walk from s to itself.

Then, we repeatedly performs a *relaxation* operation defined as follows: if (x,y) is an edge such that $dist[y] > dist[x] + c(x,y)$, then we set:

- $dist[y] = dist[x] + c(x,y)$
- $prev[y] = x$

The idea of the relaxation operation is that, if we realize that we have a better walk leading to y by using (x,y) as its last edge, compared to what we know so far, we update our knowledge.

The algorithm

Input:

G : directed graph with costs
s, t : two vertices

Output:

dist : a map that associates, to each accessible vertex, the cost of the minimum
cost walk from s to it
prev : a map that maps each accessible vertex to its predecessor on a path from s to it

Algorithm:

```
for x in X do
    dist[x] = ∞
end for
dist[s] = 0
changed = true
while changed do
    changed = false
    for (x,y) in E do
        if dist[y] > dist[x] + c(x,y) then
            dist[y] = dist[x] + c(x, y)
            prev[y] = x
            changed = true
        end if
    end for
end while
```

Proof of correctness

The proof is in three parts:

- at each stage, $dist$ and $prev$ correspond to existing walks (this comes immediately from how the relaxation operation works;
- the algorithm finishes;
- when the algorithm finishes, $dist[x] = d(s,x)$ for all vertices x .

For the last two parts, we notice that, at iteration k , we have that $dist[x] \leq w_{k,x}$ (see the [Bellman's dynamic programming algorithm](#)). This makes the Bellman-Ford finish in at most $n-1$ iterations and end with the correct distances.