

Graph algorithms - Dijkstra's algorithm

Problem

Given a graph with non-negative costs and two vertices s and t , find a minimum cost walk from s to t .

Idea

Dijkstra's algorithm still relies on Bellman's optimality principle; however, it computes distances from the starting vertex in increasing order of the distances. This way, the distance from start to a given vertex doesn't have to be recomputed after the vertex is processed.

This way, Dijkstra's algorithm looks a bit like the breadth-first traversal; however, the queue is replaced by a priority queue where the top vertex is the closest to the starting vertex.

The algorithm

```
Input:
  G : directed graph with costs
  s, t : two vertices
Output:
  dist : a map that associates, to each accessible vertex, the cost of the minimum
         cost walk from s to it
  prev : a map that maps each accessible vertex to its predecessor on a path from s to it
Algorithm:
  PriorityQueue q
  Dictionary prev
  Dictionary dist
  q.enqueue(s, 0)           // second argument is priority
  dist[s] = 0
  found = false
  while not q.isEmpty() and not found do
    x = q.dequeue()         // dequeues the element with minimum value of priority
    for y in Nout(x) do
      if y not in dist.keys() or dist[x] + cost(x,y) < dist[y] then
        dist[y] = dist[x] + cost(x, y)
        q.enqueue(y, dist[y])
        prev[y] = x
      end if
    end for
    if x == t then
      found = true
    endif
  end while
```

- If all costs are non-negative, the algorithm above doesn't put a vertex into the priority queue once it was extracted and processed (see proof below).
- If there are negative costs, but no negative cost cycles, then a vertex may be processed multiple times. However, if we eliminate the exit on dequeueing the target vertex, the algorithm finishes after a finite number of steps and the result is correct.
- If there is a negative cost cycle accessible from the starting vertex, then the algoritm can end with an incorrect result or it can run forever.

Proof of correctness (for non-negative costs)

Non-negative costs case


We claim that, when a vertex is dequeued from the priority queue, its `dist` is equal to the cost of the minimum cost walk from the start to it.

Suppose the contrary. Let x be the first vertex for which the above statement is false. So, we have that $dist[x]$ is strictly smaller than the cost of the minimum cost walk from s to x .

Let S be the set of vertices that were in the priority queue and have already been dequeued from it when x gets dequeued ($x \notin S$). On the best walk from s to x the vertex just before x cannot be in S , otherwise $dist[x]$ would have been correctly computed when that vertex was dequeued.

So, let (y,z) be the first edge on the minimum cost walk from s to x that exists S .

In the image below, the upper walk is the minimum cost walk, and the lower one is the one found by the algorithm, and implied by the values of $dist$ and $prev$.



However, since x is at the top of the priority queue and not z , we have that $cost(s,...,y,z) \geq cost(s,...,x)$ and, since all edges have non-negative costs, $cost(z,...,x) \geq 0$. Therefore, the bottom walk, found by the algorithm, cannot have a larger cost than the minimum cost walk, which prove our claim.

The case of negative costs